



# The stupid philosophers.....

## ▼ wtf r threads? an explanation for stupid people...

### ▼ analogy



- If a **process** = your apartment, a **thread** would be like *(trigger warning)* your *flatmate*.....
- All flatmates (**threads**) in the same apartment (**process**) share the same utilities (**memory**) - fridge, bathroom, TV, even your *leftovers* in the fridge if you don't put a note on it (**lock a mutex**).
- They like to do their own things *independently*: One is *cooking*, another one is *showering*, one is *watching* TV in the living room. But sometimes they *clash*:
  - You're trying to wash your underwear, but the washing machine is really cheap and shitty (*as they tend to be in shared houses*), so it doesn't automatically *lock the door*.
  - Your flatmate on the other hand, is extremely stupid (*as they tend to be in shared houses*) and opens the machine in the middle of the

wash - *causing the entire room to flood!* This is called a **race condition**. *And this is where we need a mutex....*

## ▼ and what r mu..... mutexes?

- A **mutex** = a toilet stall lock.... you lock it while you use the stuff thats inside... and unlock when you're done.
- Only one **thread** can **lock** and access the **protected data** at a time - *we don't want to poop with our friends right next to us trying to pee in the same toilet...*
- Prevents **race conditions** → two buddies (**threads**) trying to sit & shit in the toilet (**edit data**) at the same time and breaking it.

## **Mutexes in Philosophers**

|  Part |  Why It Needs a Mutex |
|--|--|
| <code>fork→mutex</code>  | Ensures only one philosopher can use a fork at a time.   |
| <code>write_mutex</code>   | Avoids messy output by letting only one message print at a time.                                       |
| <code>philo→mutex</code>   | Protects each philo's meal info ( <i>shared with monitor</i> )   |
| <code>table→mutex</code>   | Protects shared flags like <code>end_simulation</code> .   |

## ▼ a cringe SHINee analogy....

- SHINee members = **philosophers (that run threads)**
- Microphones = **forks (with mutexes)**
- They must grab **two mics** to rehearse 🎵
- They all live in the same dorm (shared memory), so coordination is fast
- Manager = **Monitor (another thread)**
- They all have a meal log (**philo→last\_meal\_time & philo→meals\_eaten**)
- The manager (**monitor**) checks those diaries (**data**) to make sure they don't miss lunch & hit their protein goals.

- To avoid smudging (**race conditions**), he only holds the diary to read while no one is writing (**mutex**)

- aka... we must put a safety lock (**mutex**) to make sure *no one can open the washing machine* (**access the shared data**) while its being used (**edited**). Or, if your flatmates are trustworthy enough, a warning sign...

## ▼ so?

- Philosophers (**structs that are used to run threads**) try to **lock** two forks (**mutexes**).
- If both forks are free, they **eat** for a set amount of time (**argv[3]**); otherwise they wait (**think**).
  - → **time\_to\_think** =  $(time\_to\_die - time\_to\_eat - time\_to\_sleep) / 2$ .
    - ./philos 5 800 200 200 → **time\_to\_think** =  $(800 - 200 - 200) / 2 = 200$ .
- After each meal, they need to sleep for a set amount of time (**argv[4]**).
- The **monitor thread** watches their `last_meal_time` using **mutexes** to avoid reading while it's being changed.

## ▼ thread functions

### ▼ pthread\_create

▼ `int pthread_create(pthread_t *thread, const pthread_attr_t *attr, void *(*start_routine)(void *), void *arg);`

Starts a new thread. duh. what else did you think.

### ▼ how to use `nr`

1. **int** - returns an error code if  $\neq 0$
2. **\*thread** - a pointer to where to store the *thread ID*
3. **\*attr** (thread attributes) - a pointer to `pthread_attr_t` object or NULL for default (enough for philosophers)

4. **\*( \*start\_routine)(void \*)** - the function that the thread will run. it must take a void\* argument and return a void \*.

5. **void \*arg** - the argument passed to the thread function ^

#### ▼ `pthread_join`

▼ `int pthread_join(pthread_t thread, void **retval);`

waits for a thread to finish whatever its doing. its just like, to make sure its cleaned up before the the main program exits i guess.

▼ how2use vv

1. **int** - error code if  $\neq 0$

2. **thread** - the thread you want to wait for (e.g. `philos[i].thread_id`)

3. **\*\*retval** (return value) - a pointer to capture the value that the function passed to **pthread\_create** returns (optional)

- **pthread\_join(philos[i].thread, NULL);** → "i dont give a FUCK what this stupid B-word wants to say to me LOL"

▼ **pthread\_join(philos[i].thread, &result);** → "ily honey tell me everything in detail..."

- **Status code or result**

- Example: `return (void *)EXIT_SUCCESS;` or `return (void *)&error_code;`

- **Result data**

- Example: a pointer to a struct with processed data or computation results.

- **Resource handles or objects**

- Example: a pointer to a buffer filled by the thread.

- **Signal completion**

- Threads can return a simple flag or pointer to indicate they finished successfully.

- **Example:** `return (NULL);`

- **Error info or messages**

- A thread can return a pointer to an error message string or detailed error info.

- ▼ `pthread_detach`

- ▼ `int pthread_detach(pthread_t thread);`

- Makes a thread run independently. Detaches the thread, so it cleans up its own resources when done. Use when you don't need to wait for a thread's result and want it to free its resources automatically.

- ▼ `void pthread_exit`

- ▼ `void pthread_exit(void *retval);`

- terminates the current thread (we don't need to call it if our thread has default attributes, our thread already does it! so I won't elaborate for now)

- ▼ `pthread_mutex_init`

- ▼ `int pthread_mutex_init(pthread_mutex_t *mutex, const pthread_mutexattr_t *attr);`

- Initializes a mutex to protect shared data from simultaneous access of course.

- ▼ `pthread_mutex_lock`

- ▼ `int pthread_mutex_lock(pthread_mutex_t *mutex);`

- locks the mutex, Protects a critical section by allowing only one thread at a time.

- ▼ `pthread_mutex_unlock`

- ▼ `int pthread_mutex_unlock(pthread_mutex_t *mutex);`

- Unlocks the mutex. Allows others to access the protected resource.

- ▼ `pthread_mutex_destroy`

▼ `int pthread_mutex_destroy(pthread_mutex_t *mutex);`

Cleans up the mutex, releases resources used by the mutex after you're done.

## ▼ the code

### ▼ library

```
#ifndef PHILO_H
# define PHILO_H

// external libraries      what we use them for:
# include <stdio.h>        // printf
# include <stdlib.h>       // malloc, free
# include <unistd.h>       // usleep
# include <sys/time.h>     // gettimeofday & struct timeval
# include <pthread.h>      // pthread_create/join, pthread_mutex_init/lock/unlock
# include <stdbool.h>      // true/false
# include <limits.h>       // INT_MAX
# include <errno.h>        // EINVAL, ENOMEM, EAGAIN, EPERM, ESRCH, EIO

/* declaring variable types */
typedef pthread_mutex_t t_mutex;
typedef struct s_table t_table; // bc we want philo in table & table in philo!

/* declaring possible operations to pass into pthread/mutex functions */
typedef enum e_opcode
{
    LOCK,
    UNLOCK,
    INIT,
    DESTROY,
    CREATE,
    JOIN,
    DETACH
}
```

```

} t_opcode;

/* declaring time units to pass into gettimeofday function */
typedef enum e_time
{
    SECOND,
    MILLISECOND,
    MICROSECOND
} t_time;

/* declaring possible actions/ states of philosophers to pass into write function */
typedef enum e_status
{
    EAT,
    SLEEP,
    THINK,
    RIGHT_FORK,
    LEFT_FORK,
    DIE,
    FULL
} t_status;

/* each fork contains: */
typedef struct s_fork
{
    t_mutex mutex;           // can only be accessed by 1 philo at a time
    int id;                  // each fork is shared by 2 philos
} t_fork;

/* each philosopher contains: */
typedef struct s_philo
{
    int id;                  // table→philos[i + 1]
    long meals_eaten;        // increases after eating
    bool full;               // if (meals_eaten == table→max_meals)
    long last_meal_time;     // time passed from last meal

```

```

    t_fork *left_fork;    // each philo needs 2 forks to eat
    t_fork *right_fork;
    pthread_t  thread_id;    // each philo runs a thread
    t_table *table;        // pointer to table data
    t_mutex  mutex;        // only one thread edits data at the same time
} t_philo;

/* the table contains : */
typedef struct  s_table
{
    long philo_nbr;        // total number of philos (argv[1])
    long time_to_die;      // how long each philo can go without eating
    long time_to_eat;      // how long it takes to eat
    long time_to_sleep;    // how long it takes to sleep
    long meals_nbr;        // argv[5] | FLAG if -1
    long start_simulation; // start time of the simulation
    bool end_simulation;   // true when a philo dies or all are full
    bool all_threads_ready; // true when all philo threads initialized & created
    long threads_running;  // how many threads running
    pthread_t  monitor;    // checks each philos meals_eaten and last_meal_eaten
    t_mutex  mutex;        // multiple threads have access so we need to lock
    t_mutex  write_mutex;  // used to synchronize writing output
    t_fork *forks;        // pointer to all forks
    t_philo *philos;      // pointer to all philos
} t_table;

/* wrapper.c */
void *handle_malloc(size_t bytes);
void handle_mutex(t_mutex *mutex, t_opcode opcode);
void handle_threads(pthread_t *thread, void *(*foo)(void *), void *data, t_opcode opcode);
/* error.c */
void error(const char *msg);
void thread_error(int status, t_opcode opcode);
void mutex_error(int status, t_opcode opcode);
/* parse.c */
void  parse_input(t_table *table, char **argv);

```



```

/* init.c */
void data_init(t_table *table);
/* dinner.c */
void start_dinner(t_table *table);
void think(t_philo *philo, bool pre_sim);
/* utils.c */
void clean(t_table *table);
long  gettime(t_time time);
void  precise_usleep(long usec, t_table *table);
/* set_get.c */
void set_bool(t_mutex *mutex, bool *dst, bool value);
void set_long(t_mutex *mutex, long *dst, long value);
long get_long(t_mutex *mutex, long *src);
bool get_bool(t_mutex *mutex, bool *src);
bool end_simulation(t_table *table);
/* sync.c */
void wait_threads(t_table *table);
bool  all_threads_running(t_mutex *mutex, long *threads, long philo_nbr);
void  increase_long(t_mutex *mutex, long *value);
void  desync_philos(t_philo *philo);
/* write.c */
void  write_status(t_status status, t_philo *philo);
/* monitor.c */
void  *monitor_dinner(void *data);
bool sim_fin(t_table *table);

#endif

```

## ▼ MAIN

```

#include "philo.h"

/* USAGE: ./philo 5(philo_nbr) 800(time_to_die) 200(time_to_eat) 200(time.
int  main(int params, char  **argv)
{
    t_table table;

```

```

if (params == 5 || params == 6) // correct input
{
    parse_input(&table, argv); // we check if its valid
    data_init(&table);         // we assign data to table, philo & fork
    philo_dinner(&table);      // here's where the magic happens
    clean(&table);             // we destroy mutexes & free memory
}
else
{
    error("Error: invalid number of arguments.\nUSAGE: ./philo number_of_philosophers");
}
return (0);
} // the rest is self explanatory, right?

```

## ▼ Checking the input

*(parse\_input)*

```

#include "philo.h"

/* functions to check for spaces & digits */
static bool is_space(char c)
{
    return ((c >= 9 && c <= 13) || c == 32);
}

static bool is_digit(char c)
{
    return (c >= '0' && c <= '9');
}

/* function to check if the input is a positive value */
static const char *valid_input(const char *str)
{
    int len;

```

```

const char *num;

len = 0;
while (is_space(*str))
    ++str;
if (*str == '+')
    ++str;
else if (*str == '-')
    error("Input Invalid: values must be positive.");
if (!is_digit(*str))
    error("Input Invalid: value must be a digit");
num = str;
while(is_digit(*str++))
    ++len;
if (len > 10)
    error("Input Invalid: value must not be larger than INT_MAX");
return (num);
}

/* atol to protect against overflow */
static long ft_atol(const char *str)
{
    long num;

    num = 0;
    str = valid_input(str);
    while (is_digit(*str))
        num = (num * 10) + (*str++ - '0');
    if (num > INT_MAX)
        error("Input Invalid: value must not be larger than INT_MAX");
    return (num);
}

/* function to check user input & write it into table struct */
void parse_input(t_table *table, char **argv)
{

```

```

table→philo_nbr = ft_atol(argv[1]);
table→time_to_die = ft_atol(argv[2]) * 1000; // time in microseconds (
table→time_to_eat = ft_atol(argv[3]) * 1000;
table→time_to_sleep = ft_atol(argv[4]) * 1000;
if (argv[5])
    table→max_meals = ft_atol(argv[5]);
else
    table→max_meals = -1;
if (table→time_to_die < 60000 || table→time_to_sleep < 60000 || tabl
    printf("Warning: 🌙 Heads up! Tiny time values (<60ms) can make
}

```

**Why do we convert to microseconds (\* 1e3/ \* 1000) ?** - Your whole philosopher program checks time using microseconds (because that's how `gettimeofday()` or `usleep()` works). So you want **everything to be in the same unit**, like everyone dancing to the same beat 🕺

## ▼ Preparing table & guests

(*data\_i nit*)

```

#include "philo.h"

/* function that fills the fork structs */
static void assign_forks(t_philo *philo, t_fork *forks, int pos)
{
    int nbr;

    nbr = philo→table→philo_nbr;
    philo→right_fork = &forks[(pos + 1) % nbr]; //wraps back to 1st fork if
    philo→left_fork = &forks[pos];
    if (philo→id % 2 == 0) // if ID is even...
    {
        philo→right_fork = &forks[pos]; //opposite assignment
        philo→left_fork = &forks[(pos + 1) % nbr];
    } // to prevent deadlock→ make philos grap 'other' fork first
}

```

```

/* function that fills the philo struct with data */
static void philo_init(t_table *table)
{
    int i;
    t_philo *philo; // to make code more readable

    i = 0;
    while (i < table->philo_nbr)
    {
        philo = table->philos + i; //assigning ptr to ptr of 'i'th philo of table
        philo->id = i + 1; //bc we dont want the first philo to be called '0'
        philo->full = false; // same as table->philos[i].full etc etc
        philo->meals_eaten = 0;
        philo->table = table; // assign table ptr so we can access table from
        handle_mutex(&philo->mutex, INIT); //bc handle_mutex expects pc
        assign_forks(philo, table->forks, i);
        i++;
    }
}

/* function that fills the table struct with data before starting the simulat
void data_init(t_table *table)
{
    int i;

    i = 0;
    table->end_simulation = false;
    table->all_threads_ready = false;
    table->threads_running = 0;
    table->philos = handle_malloc(sizeof(t_philo) * table->philo_nbr);
    handle_mutex(&table->mutex, INIT);
    handle_mutex(&table->write_mutex, INIT);
    table->forks = handle_malloc(sizeof(t_fork) * table->philo_nbr);
    while (i < table->philo_nbr)
    {

```

```

    handle_mutex(&table→forks[i].mutex, INIT);
    table→forks[i].id = i;
    i++;
} // create all forks before assigning them to philos
philo_init(table);
}

```

**why do we ask every 2nd philosopher to grab the 'other fork' first?** If everyone grabs the same side first, no one will start eating, because everyone will have one fork in their hand. That would be a deadlock.

**different ways to pass pointers:**

handle\_mutex(&table→forks[i].mutex, INIT) is the same as

```
t_fork *fork;
```

```
fork = table→fork + i;
```

```
handle_mutex(&fork→mutex, INIT);
```

## ▼ The actual dinner

*(philo<sub>d</sub>inner)*

```

void philo_dinner(t_table *table)
{
    int i;
    i = 0
    // 1. if no meals are to be eaten, skip everything
    if (table→meals_nbr == 0)
        return ;
    // 2. edge case: only 1 philo
    else if (table→philo_nbr == 1)
    {
        handle_threads(&table→philos[0].thread_id, lonely_philo, &table→p
    }
    // 3. create philo_nbr of threads
    else
    {

```

```

        while (i < table->philo_nbr)
        {
            handle_threads(&table->philos[i].thread_id, dinner_sim, &table->
                i++);
        }
    }

    // 4. create monitor thread to check if philos are dead or full
    handle_threads(&table->monitor, monitor_dinner, table, CREATE);
    // store simulation start time
    table->start_simulation = gettimeofday(MILLISECOND);
    // signal that threads may begin
    set_bool(&table->mutex, &table->all_threads_ready, true);
    i = 0;
    while (i < table->philo_nbr) // wait for philo threads to finish
        handle_threads(&table->philos[i++].thread_id, NULL, NULL, JOIN);
    set_bool(&table->mutex, &table->end_simulation, true); // all joined? s
    handle_threads(&table->monitor, NULL, NULL, JOIN); // join monitor t
}

```

▼ or **philo\_dinner: step by step** or

1. Early return if no meals needed (eeaaaasy)
2. Special case for 1 philo

▼ **lonely\_philo** hr

run\_dinner calls handle\_threads, handle\_threads calls

pthread\_create: `pthread_create(&table->philos[0].thread_id, NULL, lonely_philo, &table->philos[0]);`

telling the thread of philo[0] to run:

```

static void  *lonely_philo(void *arg) //void *arg is &table->ph
{
    t_philo *philo;

    // 1. cast void* arg to convert it back to usable philo struct
    philo = (t_philo *)arg;

```

```

// 2. wait until all_threads_ready = true
wait_threads(philo→table);
// 3. setting last_meal_time to now, simulation starts with a f
set_long(&philo→mutex, &philo→last_meal_time, gettime(
// 4. increase number of threads currently running for monit
increase_long(&philo→table→mutex, &philo→table→threa
// 5. simulate philo to pick up the only fork :,(
write_status(RIGHT_FORK, philo);
// 6. wait until he dies :(
while (!sim_fin(philo→table))
    usleep(philo→table→time_to_die);
// 7. When the simulation has ended, announce the philo's d
write_status(DIE, philo);
// 8. End the thread function by returning NULL (standard fo
return(NULL);
}

```

### 3. Create thread for each philo

#### ▼ **dinner\_sim** CUI

*run\_dinner* calls *handle\_threads*, *handle\_threads* calls *pthread\_create*:

```
pthread_create(&table→philos[i].thread_id, NULL, dinner_sim, &table-
>philos[i]);
```

telling philo[i] to run:

```

static void *dinner_sim(void *arg) // arg is &table→philo[i]
{
    t_philo *philo;
// 1. converting void * back to a philo pointer for the thread t
    philo = (t_philo *)arg;
// 2. wait until all_threads_ready == true
    wait_threads(philo→table);
// 3. set time of last meal to now, full bellies for a fair start
    set_long(&philo→mutex, &philo→last_meal_time, gettime(
// 4. let monitor know thread is running

```



```

        increase_long(&philo→table→mutex, &philo→table→threads);
// 5. desync one before starting so that no deadlocks happen
        desync_philos(philo);
// 6. run simulation loop until its over
        while (!sim_fin(philo→table))
        {
            if (philo→full)
                break ; // a. stop if full
            eat(philo); // b. take forks, set time, update meals, write status
            write_status(SLEEP, philo); // c. write sleep status
            precise_usleep(philo→table→time_to_sleep, philo→table); // d. sleep
            think(philo, false); // e. write think status and think if the
        }
// 7. exit with return NULL (standard for pthreads)
        return (NULL);
    }
}

```

#### ▼ eat

```

static void eat(t_philo *philo)
{
    handle_mutex(&philo→right_fork→mutex, LOCK); // lock right fork
    write_status(RIGHT_FORK, philo); // announce philo picked up right fork
    handle_mutex(&philo→left_fork→mutex, LOCK); // lock left fork
    write_status(LEFT_FORK, philo); // announce philo picked up left fork
    // set last_meal_time to current time
    set_long(&philo→mutex, &philo→last_meal_time, get_time());
    philo→meals_eaten++; // increase meal counter
    write_status(EAT, philo); // announce philo eating
    precise_usleep(philo→table→time_to_eat, philo→table); // sleep
    if (philo→table→max_meals > 0 && philo→meals_eaten == philo→table→max_meals)
    {
        set_bool(&philo→mutex, &philo→full, true);
        write_status(FULL, philo); // announce that philo is full
    }
}

```

```

    handle_mutex(&philo→right_fork→mutex, UNLOCK); /
    handle_mutex(&philo→left_fork→mutex, UNLOCK); //
}

```

### ▼ think

```

void think(t_philo *philo, bool pre_sim)
{
    long t_eat;
    long t_sleep;
    long t_think;

    // we only announce the philo thinking if the simulation s
    if (!pre_sim)
        write_status(THINK, philo);
    // if philo_nbr is even, we dont need thinking delays
    if (philo→table→philo_nbr % 2 == 0)
        return ;
    t_eat = philo→table→time_to_eat;
    t_sleep = philo→table→time_to_sleep;
    // calculate how long the philosopher should think
    t_think = t_eat * 2 - t_sleep;
    if (t_think < 0)
        t_think = 0; // negative thinking time = no thinking ti
    // think for 42% of thinking time, leave the rest for other
    precise_usleep(t_think * 0.42, philo→table);
}

```

### ▼ What about the other 58% of the thinking time?

The rest of the thinking time happens **naturally** as the philosopher progresses through the rest of the loop (or simulation). Here's how:

- The philosopher **sleeps** for 42% of the thinking time here, giving some initial stagger or pause.

- The **remaining time** is spent in other actions like trying to pick forks, eating, or in the main loop itself, which also takes some time.

So, the 42% sleep here is a way to **partially delay** the philosopher, just enough to desynchronize them without making them waste *all* their thinking time just waiting.

#### ▼ Why don't philosophers think in all cases?

When there's an even number of philosophers, they already take turns naturally. Half will pick up forks first, the other half will wait — no need for extra delay  
 - Desynchronization happens automatically due to the way threads are created.

#### ▼ helper functions

**wait\_threads** → simple spinlock function to check bool status every 500 microseconds.

```
void    wait_threads(t_table *table)
{
    while (!get_bool(&table->mutex, &table->all_threads_ready))
        usleep(500);
}
```

**increase\_long** → locks mutex before changing a given value, unlocks when done.

```
void    increase_long(t_mutex *mutex, long *value)
{
    handle_mutex(mutex, LOCK);
    (*value)++;
    handle_mutex(mutex, UNLOCK);
}
```

**desync\_philos** → gives certain philosophers waiting time or thinking breaks based on philo\_nbr

```

void desync_philos(t_philo *philo)
{
    // case 1: even number of philos
    if (philo->table->philo_nbr % 2 == 0)
    {
        // philo's id is even? wait until odd id philos grab the
        if (philo->id % 2 == 0)
            precise_usleep(30000, philo->table);
    }
    // case 2: odd number of philos
    else
    {
        // philo's id is odd? sleep while even id philos gra
        if (philo->id % 2)
            think(philo, true);
    }
}

```

`desync_philos` helps **stagger the start** of philosophers a little bit, so they **don't all try to grab forks at the exact same time** (which could cause chaos and deadlocks).

▼ but wait.. what's **precise\_usleep**?

`usleep()` is **not precise** enough for sensitive timing — it often **oversleeps** because:

- the OS may delay the wakeup,
- it's not real-time accurate (especially for small delays like 1–10 ms).

In the **Philosophers project**, we need **tight control** over timing — like waking up **as soon as** a philosopher has died or finished eating. That's why we implement a **more accurate sleep** manually.

```

void precise_usleep(long time_to_wait, t_table *table)
{
    long target_time;
    long time_left;

    // Save the target time in microseconds
    target_time = gettimeofday(MICROSECOND) + time_to_w

    // Loop until the total desired time has passed
    while (gettimeofday(MICROSECOND) < target_time)
    {
        // Stop early if simulation is finished
        if (sim_fin(table))
            break ;

        // How much more time do we still need to wait
        time_left = target_time - (gettimeofday(MICROSECON

        // If there's more than 1000 microseconds (1 mi
        if (time_left > 1000)
            usleep(time_left / 2); // sleep for half of the re
        else
        {
            // If under 1 ms left, wait by checking the time
            while (gettimeofday(MICROSECOND) < target_time
                ; // do nothing, just wait
        }
    }
}

```

#### 4. Create monitor thread

##### ▼ **monitor\_dinner** dr (to-do)

philo\_dinner calls handle\_threads, handle\_threads calls  
pthread\_create:

---

```
pthread_create(&table->monitor, NULL, monitor_dinner, table);
```

telling the monitor thread to run:

```
void *monitor_dinner(void *arg) // arg is table
{
    t_table *table;
    int i;
    // 1. Convert back void * into struct ptr using a cast
    table = (t_table *)arg;
    // 2. Wait until all threads are fully running
    while (!all_threads_running(&table->mutex, &table->threads))
        usleep(500);
    // 3. Keep monitoring until the simulation ends
    while (!sim_fin(table))
    {
        i = 0;
        // a. loop through each philo while the simulation is running
        while (i < table->philo_nbr && !sim_fin(table))
        {
            // b. if philo died, announce death & signal end of simulation
            if (philo_died(&table->philos[i]))
            {
                write_status(DIE, &table->philos[i]);
                set_bool(&table->mutex, &table->end_simulation, 1);
                break ;
            }
            i++;
        }
        usleep(100); // small sleep to reduce CPU stress
    }
    // 4. End the thread function by returning NULL
    return (NULL);
}
```

#### ▼ philo\_died

```

static bool philo_died(t_philo *philo)
{
    long   since_last_meal;
    long   ms_to_die;

    // If philo is full, he's alive
    if (get_bool(&philo→mutex, &philo→full))
        return (false);
    // Calculate how much time passed since last meal
    since_last_meal = gettimeofday(MILLISECOND) - get_long
    // Convert to milliseconds
    ms_to_die = philo→table→time_to_die / 1000;
    // if time_to_die has passed without meals, philo died :(
    if (since_last_meal > ms_to_die)
        return (true);
    // else he's alive :)
    return (false);
}

```

▼ Why not just use microseconds everywhere?

- Using **milliseconds** reduces the size of numbers and can be easier to work with for human-scale timing like seconds and milliseconds.
- Milliseconds are often precise enough for timing philosophers' actions.
- It keeps calculations simpler and more readable.

5. Start the simulation time and signal that threads may begin

**&table→all\_threads\_ready = true**

6. Wait for all philo threads to finish

**pthread\_join(table→philos[i].thread\_id, NULL)**

7. Mark the simulation as done

**&table→end\_simulation = true**

## 8. Join the monitor thread last

### ▼ **join\_thread explained in detail** via hpc

#### ▼ 🧵 What does it mean to "join" a thread?

Think of each philosopher like a SHINee member doing a solo stage 🎤✨.

When the show ends, the manager (your `start_dinner` function) waits behind the scenes and **doesn't leave** until each member finishes their performance.

**Joining** a thread means:

"Dear thread... I'll wait for you patiently until you're completely done before I move on."

In C, `pthread_join(thread_id, NULL)` means:

"Dear thread... I'll wait for you patiently until you're completely done before I move on."

- Pause the main thread.
- Wait for the thread with that `thread_id` to finish.
- Don't go to the next line until it's really done.

#### ▼ 🌻 What happens when we do this:

```
while (i < table->philo_nbr)
    handle_threads(&table->philos[i++].thread_id, NULL,
        NULL, JOIN);
```

This means:

1. You're going through each philosopher's thread, one by one.
2. You tell the main thread: "Hey, please wait here until **this** philosopher is finished with their simulation (they're full or died)."



3. Once a thread is done, you move to the next philosopher.
4. Only after **all philosopher threads** are done, the loop finishes.

▼ 🧠 Why do we do it?

- 💡 If you don't join threads, your program could **end too early** while philosopher threads are still running.
- 💥 It might crash or cause weird behavior like memory leaks or lost output.
- ✅ Joining makes sure everything is **clean, synchronized, and finished** before you move on.

▼ 🐾 Visual example

Let's say we have 3 philosophers, they each eat in their own thread.

```
pthread_join(philo1_thread);  
pthread_join(philo2_thread);  
pthread_join(philo3_thread);
```

It's like you're gently checking:

- 🧑 "Philo 1, are you done?"
- (waits)
- ✅ "Okay, thank you. Philo 2, are you done?"
- (waits)
- ✅ "Great job. Philo 3, are you done?"
- (waits)
- ✅ "Everyone is done. Now we can clean up and say goodbye."

▼ 🌈 And after joining?

After we've joined all philosopher threads, we can safely:

1. Signal the simulation is done  
(table→all\_threads\_ready = true);
2. Join the monitor thread (pthread\_join(\*thread, NULL));
3. Free memory, destroy mutexes  
(pthread\_mutex\_destroy(mutex));  
▼ clean(&table);

```
void clean (t_table *table)
{
    t_philo *philos;
    int i;

    i = 0;
    while (i < table→philo_nbr)
    {
        philos = table→philos + i;
        handle_mutex(&philos→mutex, DESTROY);
        i++;
    }
    handle_mutex(&table→write_mutex, DESTROY);
    handle_mutex(&table→mutex, DESTROY);
    free(table→forks);
    free(table→philos);
}
```

4. Exit the program peacefully ☁️💤 (return (0));

▼ **cleaning up** 🗑️  
(*clean*)

no further explanation needed.

```
void clean (t_table *table)
{
    int i;

    i = 0;
    // destroy each philosopher's mutex
    while (i < table->philo_nbr)
    {
        handle_mutex(&table->philos[i].mutex, DESTROY);
        handle_mutex(&table->forks[i].mutex, DESTROY);
        i++;
    }
    // destroy the mutex used for writing output
    handle_mutex(&table->write_mutex, DESTROY);
    // destroy the table mutex
    handle_mutex(&table->mutex, DESTROY);
    // free all allocated memory
    free(table->forks);
    free(table->philos);
}
```